

# segmentation\_lab

April 30, 2019

## 1 Semantic Segmentation Lab

In this lab, you will build a deep learning network that locates a particular human target within an image. The premise is that a quadcopter (simulated) is searching for a target, and then will follow the target once found. It's not enough to simply say the target is present in the image in this case, but rather to know *where* in the image the target is, so that the copter can adjust its direction in order to follow.

Consequently, an image classification network is not enough to solve the problem. Instead, a semantic segmentation network is needed so that the target can be specifically located within the image.

You can click on any of the following to quickly jump to that part of this notebook: 1. Section ?? 2. Section ?? 3. Section ?? 4. Section [1.4](#) 5. Section [1.5](#) 6. Section [1.6](#)

### 1.1 Data Collection

We have provided you with the dataset for this lab. If you haven't already downloaded the training and validation datasets, you can check out the README for this lab's repo for instructions as well.

```
In [1]: import os
import glob
import sys
import tensorflow as tf

from scipy import misc
import numpy as np

from tensorflow.contrib.keras.python import keras
from tensorflow.contrib.keras.python.keras import layers, models

from tensorflow import image

from utils import scoring_utils
from utils.separable_conv2d import SeparableConv2DKeras, BilinearUpSampling2D
from utils import data_iterator
from utils import plotting_tools
from utils import model_tools
```

```
/opt/conda/lib/python3.6/site-packages/h5py/___init___py:36: FutureWarning: Conversion of the second argument of
from ._conv import register_converters as _register_converters
```

## 1.2 FCN Layers

In the Classroom, we discussed the different layers that constitute a fully convolutional network. The following code will introduce you to the functions that you will be using to build out your model.

### 1.2.1 Separable Convolutions

The Encoder for your FCN will essentially require separable convolution layers. Below we have implemented two functions - one which you can call upon to build out separable convolutions or regular convolutions. Each with batch normalization and with the ReLU activation function applied to the layers.

While we recommend the use of separable convolutions thanks to their advantages we covered in the Classroom, some of the helper code we will present for your model will require the use for regular convolutions. But we encourage you to try and experiment with each as well!

The following will help you create the encoder block and the final model for your architecture.

```
In [2]: def separable_conv2d_batchnorm(input_layer, filters, strides=1):
        output_layer = SeparableConv2DKeras(filters=filters, kernel_size=3, strides=strides,
                                             padding='same', activation='relu')(input_layer)

        output_layer = layers.BatchNormalization()(output_layer)
        return output_layer

    def conv2d_batchnorm(input_layer, filters, kernel_size=3, strides=1):
        output_layer = layers.Conv2D(filters=filters, kernel_size=kernel_size, strides=strides,
                                     padding='same', activation='relu')(input_layer)

        output_layer = layers.BatchNormalization()(output_layer)
        return output_layer
```

### 1.2.2 Bilinear Upsampling

The following helper function will help implement the bilinear upsampling layer. Currently, up-sampling by a factor of 2 is recommended but you can try out different factors as well. You will use this to create the decoder block later!

```
In [3]: def bilinear_upsample(input_layer):
        output_layer = BilinearUpSampling2D((2,2))(input_layer)
        return output_layer
```

## 1.3 Build the Model

In the following cells, we will cover how to build the model for the task at hand.

- We will first create an Encoder Block, where you will create a separable convolution layer using an input layer and the size(depth) of the filters as your inputs.
- Next, you will create the Decoder Block, where you will create an upsampling layer using bilinear upsampling, followed by a layer concatenation, and some separable convolution layers.
- Finally, you will combine the above two and create the model. In this step you will be able to experiment with different number of layers and filter sizes for each to build your model.

Let's cover them individually below.

### 1.3.1 Encoder Block

Below you will create a separable convolution layer using the `separable_conv2d_batchnorm()` function. The `filters` parameter defines the size or depth of the output layer. For example, 32 or 64.

```
In [4]: def encoder_block(input_layer, filters, strides=1):

        # TODO Create a separable convolution layer using the separable_conv2d_batchnorm() f
        output_layer = separable_conv2d_batchnorm(input_layer, filters=filters, strides=stri
        return output_layer
```

### 1.3.2 Decoder Block

The decoder block, as covered in the Classroom, comprises of three steps -

- A bilinear upsampling layer using the `bilinear_upsample()` function. The current recommended factor for upsampling is set to 2.
- A layer concatenation step. This step is similar to skip connections. You will concatenate the upsampled `small_ip_layer` and the `large_ip_layer`.
- Some (one or two) additional separable convolution layers to extract some more spatial information from prior layers.

```
In [5]: def decoder_block(small_ip_layer, large_ip_layer, filters):

        # TODO Upsample the small input layer using the bilinear_upsample() function.
        Upsampled_small_ip_layer = bilinear_upsample(small_ip_layer)
        # TODO Concatenate the upsampled and large input layers using layers.concatenate
        concatenated_layers = layers.concatenate([Upsampled_small_ip_layer, large_ip_layer])
        # TODO Add some number of separable convolution layers
        output_layer = encoder_block(concatenated_layers, filters=filters, strides=1)
        return output_layer
```

### 1.3.3 Model

Now that you have the encoder and decoder blocks ready, you can go ahead and build your model architecture!

There are three steps to the following: - Add encoder blocks to build out initial set of layers. This is similar to how you added regular convolutional layers in your CNN lab. - Add 1x1 Convolution layer using `conv2d_batchnorm()` function. Remember that 1x1 Convolutions require a kernel and stride of 1. - Add decoder blocks for upsampling and skip connections.

```
In [20]: def fcn_model(inputs, num_classes):

    # TODO Add Encoder Blocks.
    # Remember that with each encoder layer, the depth of your model (the number of fil
    print(inputs.get_shape())
    #Layer 1
    layer_1 = encoder_block(inputs, filters=32, strides=2)
    print(layer_1.get_shape())
    #Layer 2
    layer_2 = encoder_block(layer_1, filters=64, strides=2)
    print(layer_2.get_shape())
    # TODO Add 1x1 Convolution layer using conv2d_batchnorm().
    layer_1x1 = conv2d_batchnorm(layer_2, filters=128, kernel_size=3, strides=2)
    print(layer_1x1.get_shape())
    # TODO: Add the same number of Decoder Blocks as the number of Encoder Blocks
    decoder_1 = decoder_block(layer_1x1, layer_2, filters=64)
    print(decoder_1.get_shape())
    decoder_2 = decoder_block(decoder_1, layer_1, filters=32)
    print(decoder_2.get_shape())
    x = decoder_block(decoder_2, inputs, filters=3)
    print(x.get_shape())
    # The function returns the output layer of your model. "x" is the final layer obtain
    return layers.Conv2D(num_classes, 3, activation='softmax', padding='same')(x)
```

## 1.4 Training

The following cells will utilize the model you created and define an output layer based on the input and the number of classes. Following that you will define the hyperparameters to compile and train your model!

```
In [21]: """
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

image_hw = 128
image_shape = (image_hw, image_hw, 3)
inputs = layers.Input(image_shape)
num_classes = 3

# Call fcn_model()
output_layer = fcn_model(inputs, num_classes)

(?, 128, 128, 3)
(?, 64, 64, 32)
```

```
(?, 32, 32, 64)
(?, 16, 16, 128)
(?, 32, 32, 64)
(?, 64, 64, 32)
(?, 128, 128, 3)
```

### 1.4.1 Hyperparameters

Define and tune your hyperparameters. - **batch\_size**: number of training samples/images that get propagated through the network in a single pass. - **num\_epochs**: number of times the entire training dataset gets propagated through the network. - **steps\_per\_epoch**: number of batches of training images that go through the network in 1 epoch. We have provided you with a default value. One recommended value to try would be based on the total number of images in training dataset divided by the batch\_size. - **validation\_steps**: number of batches of validation images that go through the network in 1 epoch. This is similar to steps\_per\_epoch, except validation\_steps is for the validation dataset. We have provided you with a default value for this as well. - **workers**: maximum number of processes to spin up. This can affect your training speed and is dependent on your hardware. We have provided a recommended value to work with.

```
In [22]: learning_rate = 0.003
        batch_size = 100
        num_epochs = 20
        steps_per_epoch = 50
        validation_steps = 50
        workers = 2
```

```
In [23]: """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """

        # Define the Keras model and compile it for training
        model = models.Model(inputs=inputs, outputs=output_layer)

        model.compile(optimizer=keras.optimizers.Adam(learning_rate), loss='categorical_crossentropy')

        # Data iterators for loading the training and validation data
        train_iter = data_iterator.BatchIteratorSimple(batch_size=batch_size,
                                                       data_folder=os.path.join('..', 'data', 'train'),
                                                       image_shape=image_shape,
                                                       shift_aug=True)

        val_iter = data_iterator.BatchIteratorSimple(batch_size=batch_size,
                                                      data_folder=os.path.join('..', 'data', 'validation'),
                                                      image_shape=image_shape)

        logger_cb = plotting_tools.LoggerPlotter()
        callbacks = [logger_cb]
```

```

model.fit_generator(train_iter,
                    steps_per_epoch = steps_per_epoch, # the number of batches per epoch
                    epochs = num_epochs, # the number of epochs to train for,
                    validation_data = val_iter, # validation iterator
                    validation_steps = validation_steps, # the number of batches to validate
                    callbacks=callbacks,
                    workers = workers)

```


Epoch 1/20

```

/opt/conda/lib/python3.6/site-packages/scipy/misc/pilutil.py:482: FutureWarning: Conversion of the
if issubdtype(ts, int):
/opt/conda/lib/python3.6/site-packages/scipy/misc/pilutil.py:485: FutureWarning: Conversion of the
elif issubdtype(type(size), float):

```


49/50 [=====>.] - ETA: 0s - loss: 0.8849



50/50 [=====] - 49s - loss: 0.8771 - val\_loss: 0.5622

Epoch 2/20


49/50 [=====>.] - ETA: 0s - loss: 0.2171




50/50 [=====] - 45s - loss: 0.2147 - val\_loss: 0.1372

Epoch 3/20

49/50 [=====>.] - ETA: 0s - loss: 0.0724




50/50 [=====] - 45s - loss: 0.0721 - val\_loss: 0.1147  
Epoch 4/20  
49/50 [=====>.] - ETA: 0s - loss: 0.0581




50/50 [=====] - 45s - loss: 0.0581 - val\_loss: 0.1092  
Epoch 5/20  
49/50 [=====>.] - ETA: 0s - loss: 0.0495







50/50 [=====] - 45s - loss: 0.0494 - val\_loss: 0.1029  
 Epoch 6/20  
 49/50 [=====>.] - ETA: 0s - loss: 0.0443




50/50 [=====] - 45s - loss: 0.0443 - val\_loss: 0.0971  
Epoch 7/20  
49/50 [=====>.] - ETA: 0s - loss: 0.0410




50/50 [=====] - 45s - loss: 0.0408 - val\_loss: 0.0889  
Epoch 8/20  
49/50 [=====>.] - ETA: 0s - loss: 0.0369




50/50 [=====] - 45s - loss: 0.0369 - val\_loss: 0.0833  
 Epoch 9/20  
 49/50 [=====>.] - ETA: 0s - loss: 0.0393



50/50 [=====] - 45s - loss: 0.0391 - val\_loss: 0.0666  
Epoch 10/20  
49/50 [=====>.] - ETA: 0s - loss: 0.0344




50/50 [=====] - 45s - loss: 0.0343 - val\_loss: 0.0566  
Epoch 11/20  
49/50 [=====>.] - ETA: 0s - loss: 0.0337




50/50 [=====] - 45s - loss: 0.0338 - val\_loss: 0.0555

Epoch 12/20


49/50 [=====>.] - ETA: 0s - loss: 0.0322



50/50 [=====] - 44s - loss: 0.0321 - val\_loss: 0.0510  
Epoch 13/20  
49/50 [=====>.] - ETA: 0s - loss: 0.0298




50/50 [=====] - 45s - loss: 0.0299 - val\_loss: 0.0523  
Epoch 14/20  
49/50 [=====>.] - ETA: 0s - loss: 0.0307




```

50/50 [=====] - 45s - loss: 0.0307 - val_loss: 0.0491
Epoch 15/20
49/50 [=====>.] - ETA: 0s - loss: 0.0285
  
```




50/50 [=====] - 45s - loss: 0.0284 - val\_loss: 0.0451  
Epoch 16/20  
49/50 [=====>.] - ETA: 0s - loss: 0.0289




50/50 [=====] - 45s - loss: 0.0288 - val\_loss: 0.0498  
Epoch 17/20  
49/50 [=====>.] - ETA: 0s - loss: 0.0274







50/50 [=====] - 45s - loss: 0.0275 - val\_loss: 0.0463  
 Epoch 18/20  
 49/50 [=====>.] - ETA: 0s - loss: 0.0272



50/50 [=====] - 45s - loss: 0.0272 - val\_loss: 0.0438  
Epoch 19/20  
49/50 [=====>.] - ETA: 0s - loss: 0.0260



50/50 [=====] - 45s - loss: 0.0260 - val\_loss: 0.0477  
Epoch 20/20  
49/50 [=====>.] - ETA: 0s - loss: 0.0261



50/50 [=====] - 45s - loss: 0.0261 - val\_loss: 0.0386

Out[23]: <tensorflow.contrib.keras.python.keras.callbacks.History at 0x7fc7f5ca5438>

```
In [24]: # Save your trained model weights
weight_file_name = 'model_weights'
model_tools.save_network(model, weight_file_name)
```

## 1.5 Prediction

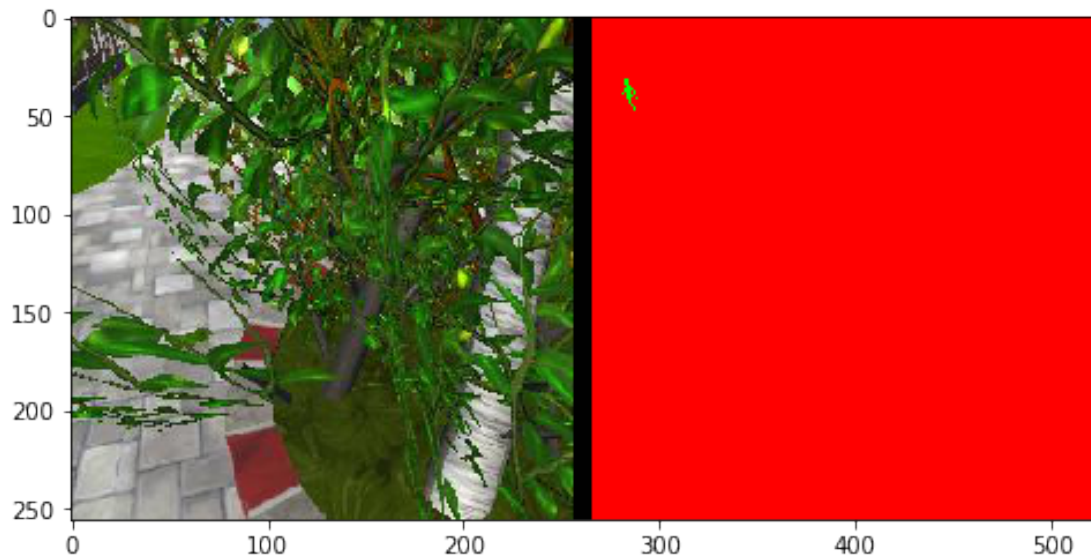
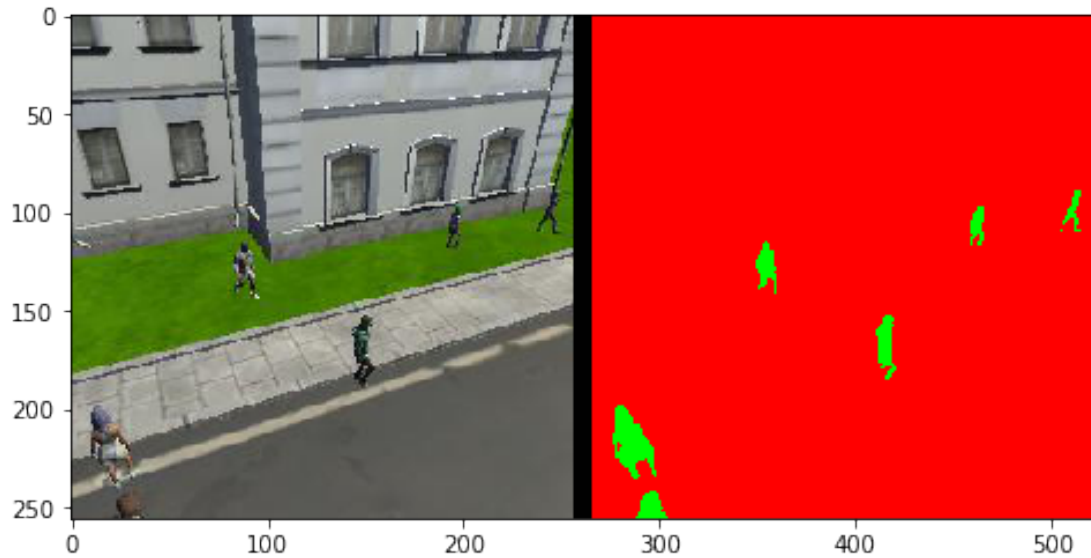
In [25]: *# If you need to load a model which you previously trained you can uncomment the code below*

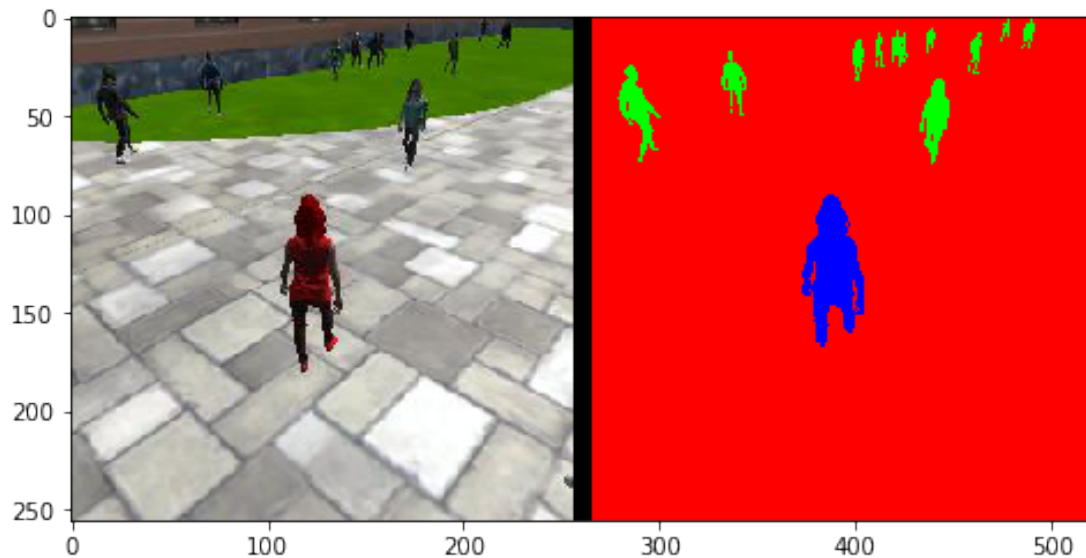
```
weight_file_name = 'model_weights'
restored_model = model_tools.load_network(weight_file_name)
```

```
In [26]: # generate predictions, save in the runs, directory.
run_number = 'run1'
validation_path, output_path = model_tools.write_predictions_grade_set(model, run_number)
```

```
/opt/conda/lib/python3.6/site-packages/scipy/misc/pilutil.py:482: FutureWarning: Conversion of the path from '...' to unicode is deprecated.
if issubdtype(ts, int):
/opt/conda/lib/python3.6/site-packages/scipy/misc/pilutil.py:485: FutureWarning: Conversion of the path from '...' to unicode is deprecated.
elif issubdtype(type(size), float):
```

```
In [27]: # take a look at predictions
# validation_path = 'validation'
im_files = plotting_tools.get_im_file_sample(run_number, validation_path)
for i in range(3):
    im_tuple = plotting_tools.load_images(im_files[i])
    plotting_tools.show_images(im_tuple)
```





## 1.6 Evaluation

Let's evaluate your model!

```
In [28]: scoring_utils.score_run(validation_path, output_path)
```

```
/opt/conda/lib/python3.6/site-packages/scipy/misc/pilutil.py:482: FutureWarning: Conversion of the path from 'validation_path' to
if issubdtype(ts, int):
/opt/conda/lib/python3.6/site-packages/scipy/misc/pilutil.py:485: FutureWarning: Conversion of the path from 'output_path' to
elif issubdtype(type(size), float):
```

```
number of validation samples intersection over the union evaluated on 1184
average intersection over union for background is 0.9911219735422065
average intersection over union for other people is 0.27819893422634084
average intersection over union for hero is 0.14170643840902072
global average intersection over union is 0.47034244872585607
```